

Start coding or [generate](#) with AI.

```
# Example: Baseline inference with a 1B instruct model (zero-shot)
!pip install unsloth transformers datasets -q
from datasets import load_dataset
from unsloth import FastLanguageModel
import torch

# Load data (take 75%-agreement subset) and split
dataset = load_dataset("takala/financial_phrasebank", "sentences_75agree")
data = dataset["train"]
data = data.train_test_split(test_size=0.2, seed=42)
train_data = data["train"]
test_data = data["test"]

# Load base model in 4-bit (QLoRA) for speed
model_name = "unsloth/llama-3.2-1B-Instruct-bnb-4bit"
model, tokenizer = FastLanguageModel.from_pretrained(model_name, max_seq_length=2048, load_in_4bit=True)
FastLanguageModel.for_inference(model) # prepare for faster generation

# Example inference on one sample
sample = test_data[0]["sentence"]
prompt = f"{{sample}} Sentiment:"
inputs = tokenizer(prompt, return_tensors="pt").to("cuda")
outs = model.generate(input_ids=inputs['input_ids'], max_new_tokens=3, temperature=0.2)
label = tokenizer.decode(outs[0], skip_special_tokens=True)
print("Sentence:", sample)
print("Predicted sentiment:", label)
```



```
_____ 52.3/52.3 kB 3.6 MB/s eta 0:00:00
_____ 311.7/311.7 kB 13.6 MB/s eta 0:00:00
_____ 491.5/491.5 kB 30.2 MB/s eta 0:00:00
_____ 511.9/511.9 kB 44.9 MB/s eta 0:00:00
_____ 184.8/184.8 kB 18.6 MB/s eta 0:00:00
_____ 117.2/117.2 MB 8.0 MB/s eta 0:00:00
_____ 61.3/61.3 MB 13.4 MB/s eta 0:00:00
_____ 130.0/130.0 kB 12.0 MB/s eta 0:00:00
_____ 213.6/213.6 kB 20.2 MB/s eta 0:00:00
```

🦄 Unsloth: Will patch your computer to enable 2x faster free finetuning.

🦄 Unsloth Zoo will now patch everything to make training faster!

README.md: 8.88k/? [00:00<00:00, 711kB/s]

financial\_phrasebank.py: 6.04k/? [00:00<00:00, 578kB/s]

The repository for takala/financial\_phrasebank contains custom code which must be executed to correctly load the dataset. You can inspect the code. You can avoid this prompt in future by passing the argument `trust\_remote\_code=True`.

Do you wish to run the custom code? [y/N] y

FinancialPhraseBank-v1.0.zip: 100% 682k/682k [00:00<00:00, 8.38MB/s]

Generating train split: 100% 3453/3453 [00:00<00:00, 42658.84 examples/s]

==((====))== Unsloth 2025.8.9: Fast Llama patching. Transformers: 4.55.4.

\\ /| Tesla T4. Num GPUs = 1. Max memory: 14.741 GB. Platform: Linux.

0^0/ \\_/ \ Torch: 2.8.0+cu126. CUDA: 7.5. CUDA Toolkit: 12.6. Triton: 3.4.0

\ \_\_\_\_/ Bfloat16 = FALSE. FA [Xformers = 0.0.32.post2. FA2 = False]

"-\_\_-" Free license: <http://github.com/unslothai/unsloth>

Unsloth: Fast downloading is enabled - ignore downloading bars which are red colored!

model.safetensors: 100% 1.03G/1.03G [00:08<00:00, 221MB/s]

generation\_config.json: 100% 234/234 [00:00<00:00, 14.4kB/s]

tokenizer\_config.json: 54.7k/? [00:00<00:00, 5.27MB/s]

tokenizer.json: 100% 17.2M/17.2M [00:00<00:00, 49.1MB/s]

special\_tokens\_map.json: 100% 454/454 [00:00<00:00, 37.5kB/s]

The attention mask is not set and cannot be inferred from input because pad token is same as eos token. As a consequence, you may observe Sentence: It estimates the operating profit to further improve from the third quarter .

Predicted sentiment: It estimates the operating profit to further improve from the third quarter. Sentiment: positive. The

```
print(test_data[3])
```



```
{'sentence': 'Our tools are specifically designed with the needs of both the business users and ICT experts in mind .', 'label': 1}
```

```

!pip install scikit-learn -q
from sklearn.metrics import accuracy_score, f1_score

# Label mapping
label_map = {"negative": 0, "neutral": 1, "positive": 2}

def normalize_pred(text):
    text = text.lower()
    if "pos" in text:
        return 2 # positive
    elif "neg" in text:
        return 0 # negative
    elif "neu" in text:
        return 1 # neutral
    else:
        return 1 # fallback to neutral

y_true = []
y_pred = []

for sample in test_data:
    sent = sample["sentence"]
    gold = sample["label"] # already int (0/1/2)

    prompt = f"Sentence: {sent}\nSentiment:"
    inputs = tokenizer(prompt, return_tensors="pt").to("cuda")

    outputs = model.generate(
        input_ids=inputs["input_ids"],
        attention_mask=inputs["attention_mask"],
        max_new_tokens=3,
        temperature=0.2,
    )
    pred_text = tokenizer.decode(outputs[0], skip_special_tokens=True)
    pred = normalize_pred(pred_text)

    y_true.append(gold)
    y_pred.append(pred)

# Compute metrics
acc = accuracy_score(y_true, y_pred)
f1 = f1_score(y_true, y_pred, average="macro")
print(f"Baseline Accuracy: {acc*100:.2f}%")
print(f"Baseline F1 (macro): {f1*100:.2f}%")

```

↗ Baseline Accuracy: 61.79%  
Baseline F1 (macro): 30.51%

```

from unsloth import FastLanguageModel

# Load base model (16-bit for LoRA training)
model, tokenizer = FastLanguageModel.from_pretrained(
    "unsloth/Llama-3.2-1B",
    max_seq_length=2048,
    load_in_4bit=False,
)

# Wrap with LoRA
model = FastLanguageModel.get_peft_model(
    model,
    r=16,
    target_modules=["q_proj", "k_proj", "v_proj", "o_proj", "gate_proj", "up_proj", "down_proj"],
    lora_alpha=16,
    lora_dropout=0.0,
    bias="none",
    use_gradient_checkpointing="unsloth",
)

model.print_trainable_parameters()

```

```

==((====))==  Unsloth 2025.8.9: Fast Llama patching. Transformers: 4.55.4.
  \ \  /|   Tesla T4. Num GPUs = 1. Max memory: 14.741 GB. Platform: Linux.
0^0/ \_/_\   Torch: 2.8.0+cu126. CUDA: 7.5. CUDA Toolkit: 12.6. Triton: 3.4.0
 \  _ _ _/   Bfloat16 = FALSE. FA [Xformers = 0.0.32.post2. FA2 = False]
  "-__-"     Free license: http://github.com/unslothai/unsloth
Unsloth: Fast downloading is enabled - ignore downloading bars which are red colored!
model.safetensors: 100%                2.47G/2.47G [00:32<00:00, 262MB/s]

generation_config.json: 100%            230/230 [00:00<00:00, 15.0kB/s]

tokenizer_config.json:   50.6k/? [00:00<00:00, 4.92MB/s]

special_tokens_map.json: 100%            459/459 [00:00<00:00, 52.5kB/s]

tokenizer.json: 100%                17.2M/17.2M [00:00<00:00, 31.5MB/s]
Unsloth 2025.8.9 patched 16 layers with 16 QKV layers, 16 O layers and 16 MLP layers.
trainable params: 11,272,192 || all params: 1,247,086,592 || trainable%: 0.9039

```

```
id2label = {0: "negative", 1: "neutral", 2: "positive"}
```

```

def formatting_func(example):
    texts = []
    sentences = example["sentence"]
    labels = example["label"]

    # Case 1: batch (lists)
    if isinstance(sentences, list):
        for sent, lab in zip(sentences, labels):
            texts.append(f"Sentence: {sent}\nSentiment: {id2label[lab]}")
    else:
        # Case 2: single sample
        texts.append(f"Sentence: {sentences}\nSentiment: {id2label[labels]}")
    return texts

```

```

from transformers import TrainingArguments
from unsloth.trainer import SFTTrainer

```

```

args = TrainingArguments(
    output_dir="outputs_lora_1b",
    num_train_epochs=3,
    per_device_train_batch_size=2,
    gradient_accumulation_steps=4,
    learning_rate=2e-4,
    fp16=True,
    logging_steps=50,
    save_strategy="no",
)

```

```

trainer = SFTTrainer(
    model=model,
    tokenizer=tokenizer,
    train_dataset=train_data,
    eval_dataset=test_data,
    formatting_func=formatting_func,
    max_seq_length=512,
    args=args,
)

```

```
trainer.train()
```

```

Unsloth: Tokenizing ["text"]: 100%                2762/2762 [00:01<00:00, 2435.79 examples/s]

Unsloth: Tokenizing ["text"]: 100%                691/691 [00:00<00:00, 2955.41 examples/s]

==(====)= Unsloth - 2x faster free finetuning | Num GPUs used = 1
  \ \ / | Num examples = 2,762 | Num Epochs = 3 | Total steps = 1,038
0^0/ \_/ \ Batch size per device = 2 | Gradient accumulation steps = 4
 \ \ / / Data Parallel GPUs = 1 | Total batch size (2 x 4 x 1) = 8
"-_____" Trainable parameters = 11,272,192 of 1,247,086,592 (0.90% trained)
/usr/local/lib/python3.12/dist-packages/notebook/notebookapp.py:191: SyntaxWarning: invalid escape sequence '\/'
|_| | '_ \ / _ \ | _/ -_)
wandb: Logging into wandb.ai. (Learn how to deploy a W&B server locally: https://wandb.me/wandb-server)
wandb: You can find your API key in your browser here: https://wandb.ai/authorize?ref=models
wandb: Paste an API key from your profile and hit enter: wandb: WARNING If you're specifying your api key in code, ensure this code is
wandb: WARNING Consider setting the WANDB_API_KEY environment variable, or running `wandb login` from the command line.
wandb: No netrc file found, creating one.
wandb: Appending key for api.wandb.ai to your netrc file: /root/.netrc
wandb: Currently logged in as: aravindyuvraj007 (aravindyuvraj007-vnrvjietofficial) to https://api.wandb.ai. Use `wandb login --relogi
Tracking run with wandb version 0.21.1
Run data is saved locally in /content/wandb/run-20250828_112757-19xaw7pe
Syncing run easy-universe-2 to Weights & Biases \(docs\)
View project at https://wandb.ai/aravindyuvraj007-vnrvjietofficial/huggingface
View run at https://wandb.ai/aravindyuvraj007-vnrvjietofficial/huggingface/runs/19xaw7pe
[ 362/1038 04:37 < 08:41, 1.30 it/s, Epoch 1.04/3]

```

| Step | Training Loss |
|------|---------------|
|------|---------------|

|     |          |
|-----|----------|
| 50  | 2.756500 |
| 100 | 2.333300 |
| 150 | 2.237900 |
| 200 | 2.260500 |
| 250 | 2.180100 |
| 300 | 2.213300 |
| 350 | 2.137000 |

[1038/1038 13:16, Epoch 3/3]

| Step | Training Loss |
|------|---------------|
|------|---------------|

|     |          |
|-----|----------|
| 50  | 2.756500 |
| 100 | 2.333300 |
| 150 | 2.237900 |
| 200 | 2.260500 |
| 250 | 2.180100 |
| 300 | 2.213300 |
| 350 | 2.137000 |
| 400 | 1.888300 |
| 450 | 1.904700 |

```
from sklearn.metrics import accuracy_score, f1_score
```

```
y_true, y_pred = [], []
```

```
id2label = {0: "negative", 1: "neutral", 2: "positive"}
label2id = {v:k for k,v in id2label.items()}
```

```
for ex in test_data:
    prompt = f"Sentence: {ex['sentence']}\nSentiment:"
    inputs = tokenizer(prompt, return_tensors="pt").to(model.device)
    outputs = model.generate(**inputs, max_new_tokens=5)
    pred = tokenizer.decode(outputs[0], skip_special_tokens=True)
```

```
# extract last token as sentiment guess
```

```
guess = None
for cand in id2label.values():
    if cand in pred.lower():
        guess = cand
        break
```

```
if guess is None:
    guess = "neutral" # fallback
```

```
y_true.append(ex["label"])
```

```

y_pred.append(label2id[guess])

# Compute metrics
acc = accuracy_score(y_true, y_pred)
f1 = f1_score(y_true, y_pred, average="macro")
print(f"Fine-tuned LoRA Accuracy: {acc*100:.2f}%")
print(f"Fine-tuned LoRA F1 (macro): {f1*100:.2f}%")

```

↗ Fine-tuned LoRA Accuracy: 90.74%  
Fine-tuned LoRA F1 (macro): 89.05%

Start coding or [generate](#) with AI.

```

# Test the fine-tuned model with a new prompt
new_prompt = "Sentence: I am not feeling good, I am feeling very not good and not good heart attack.\nSentiment:"
inputs = tokenizer(new_prompt, return_tensors="pt").to(model.device)

outputs = model.generate(**inputs, max_new_tokens=5)
generated_text = tokenizer.decode(outputs[0], skip_special_tokens=True)

print("Generated Sentiment:", generated_text)

```

↗ Generated Sentiment: Sentence: I am not feeling good, I am feeling very not good and not good heart attack.  
Sentiment: negative - negative Sentiment

Start coding or [generate](#) with AI.

## ✓ Task

Set up the code to fine-tune a language model using QLoRA with the `FastLanguageModel` library. Load the model in 4-bit precision, wrap it with LoRA, define training arguments, and initialize the `SFTTrainer`. Do not run the training.

## ✓ Load the model in 4-bit

### Subtask:

Load the base model in 4-bit precision using `FastLanguageModel.from_pretrained`.

**Reasoning:** The subtask is to load the base model in 4-bit precision. This requires using the `FastLanguageModel.from_pretrained` method with the specified parameters.

```

from unsloth import FastLanguageModel

# Load base model in 4-bit for QLoRA fine-tuning
model, tokenizer = FastLanguageModel.from_pretrained(
    "unsloth/Llama-3.2-3B-bnb-4bit",
    max_seq_length=2048,
    load_in_4bit=True,
)

peft_model = FastLanguageModel.get_peft_model(
    model,
    r=16,
    target_modules=["q_proj", "k_proj", "v_proj", "o_proj", "gate_proj", "up_proj", "down_proj"],
    lora_alpha=16,
    lora_dropout=0.0,
    bias="none",
    use_gradient_checkpointing="unsloth"
)

peft_model.print_trainable_parameters()

```

↗ ==((====))== Unsloth 2025.8.9: Fast Llama patching. Transformers: 4.55.4.  
\\ /| Tesla T4. Num GPUs = 1. Max memory: 14.741 GB. Platform: Linux.  
0^0/ \\_/ \ Torch: 2.8.0+cu126. CUDA: 7.5. CUDA Toolkit: 12.6. Triton: 3.4.0  
\\ \_\_\_\_/ Bfloat16 = FALSE. FA [Xformers = 0.0.32.post2. FA2 = False]  
"-\_\_\_\_\_" Free license: <http://github.com/unslothai/unsloth>


Unsloth: Fast downloading is enabled - ignore downloading bars which are red colored!  
Unsloth 2025.8.9 patched 28 layers with 28 QKV layers, 28 O layers and 28 MLP layers.  
trainable params: 24,313,856 || all params: 3,237,063,680 || trainable%: 0.7511

```
from transformers import TrainingArguments
from unsloth.trainer import SFTTrainer
```

```
args = TrainingArguments(
    output_dir="outputs_lora_1b",
    num_train_epochs=3,
    per_device_train_batch_size=2,
    gradient_accumulation_steps=4,
    learning_rate=2e-4,
    fp16=True,
    logging_steps=50,
    save_strategy="no",
)
```

```
trainer = SFTTrainer(
    model=peft_model,
    tokenizer=tokenizer,
    train_dataset=train_data,
    eval_dataset=test_data,
    formatting_func=formatting_func,
    max_seq_length=512,
    args=args,
)
```

```
trainer.train()
```



Unsloth: Tokenizing ["text"]: 100%

2762/2762 [00:00<00:00, 5458.54 examples/s]

Unsloth: Tokenizing ["text"]: 100%

691/691 [00:00<00:00, 4953.17 examples/s]

==(====)=

Unsloth - 2x faster free finetuning | Num GPUs used = 1

\\ / | Num examples = 2,762 | Num Epochs = 3 | Total steps = 1,038

0^0/ \\_ / \ Batch size per device = 2 | Gradient accumulation steps = 4

\\ / \ Data Parallel GPUs = 1 | Total batch size (2 x 4 x 1) = 8

"-\_\_\_-" Trainable parameters = 24,313,856 of 3,237,063,680 (0.75% trained)

[1038/1038 29:14, Epoch 3/3]

Step Training Loss

|      |          |
|------|----------|
| 50   | 2.503700 |
| 100  | 2.219600 |
| 150  | 2.118500 |
| 200  | 2.134100 |
| 250  | 2.057000 |
| 300  | 2.091300 |
| 350  | 2.022600 |
| 400  | 1.771900 |
| 450  | 1.780200 |
| 500  | 1.790400 |
| 550  | 1.720600 |
| 600  | 1.777300 |
| 650  | 1.723300 |
| 700  | 1.671700 |
| 750  | 1.438500 |
| 800  | 1.431600 |
| 850  | 1.447000 |
| 900  | 1.402600 |
| 950  | 1.383300 |
| 1000 | 1.385300 |

TrainOutput(global\_step=1038, training\_loss=1.7802981264559061, metrics={'train\_runtime': 1759.1777, 'train\_samples\_per\_second': 4.71, 'train\_steps\_per\_second': 0.59, 'total\_flos': 6197031974510592.0, 'train\_loss': 1.7802981264559061})

```
from sklearn.metrics import accuracy_score, f1_score

y_true, y_pred = [], []

id2label = {0: "negative", 1: "neutral", 2: "positive"}
label2id = {v:k for k,v in id2label.items()}

for ex in test_data:
    prompt = f"Sentence: {ex['sentence']}\nSentiment:"
    inputs = tokenizer(prompt, return_tensors="pt").to(model.device)
    outputs = peft_model.generate(**inputs, max_new_tokens=5)
    pred = tokenizer.decode(outputs[0], skip_special_tokens=True)

    # extract last token as sentiment guess
    guess = None
    for cand in id2label.values():
        if cand in pred.lower():
            guess = cand
            break
    if guess is None:
        guess = "neutral" # fallback

    y_true.append(ex["label"])
    y_pred.append(label2id[guess])

# Compute metrics
acc = accuracy_score(y_true, y_pred)
f1 = f1_score(y_true, y_pred, average="macro")
print(f"Fine-tuned LoRA Accuracy: {acc*100:.2f}%")
```

```
print(f"Fine-tuned LoRA F1 (macro): {f1*100:.2f}%")
```

↗ Fine-tuned LoRA Accuracy: 95.66%  
Fine-tuned LoRA F1 (macro): 94.98%

```
from peft import PromptTuningConfig, get_peft_model, TaskType, PromptTuningInit
```